

A case study of time-series forecasting for bike rentals

Open Source Experience 2022

Guillaume Lemaitre - November 9, 2022



About me

Guillaume Lemaître - Software Engineer

2017

2019

2021



scikit-learn consortium
@ Inria La Fondation





Core developer



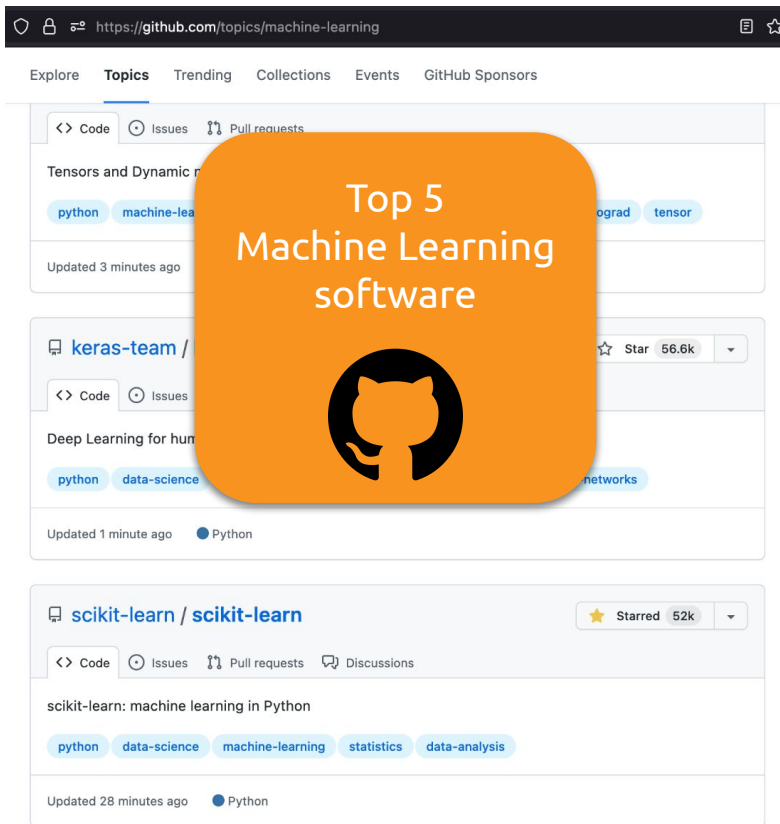
Core developer



 @glemaitre

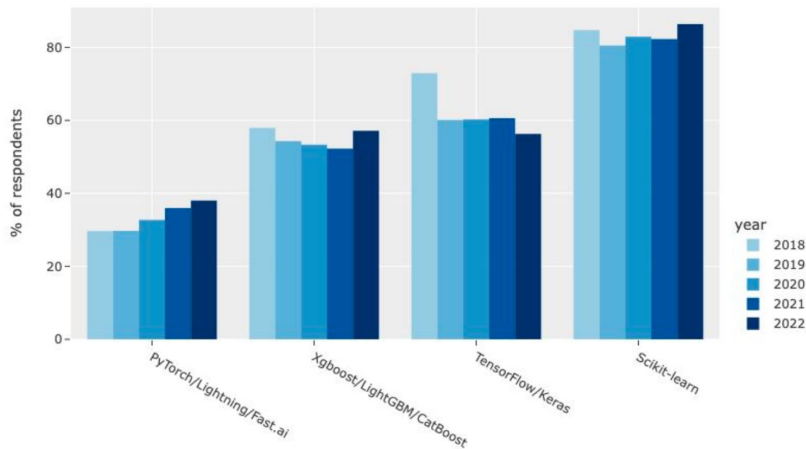
 @glemaitre@fosstodon.org

About scikit-learn



Kaggle DS & ML Survey 2022

Scikit-learn is the most popular ML framework while PyTorch has been growing steadily year-over-year



Agenda

Introduction to time-series data

Mind the evaluation

Time-series forecasting

Modeling predictive uncertainty

Beyond scikit-learn

Tackling time-series as a regression problem

The bike sharing dataset

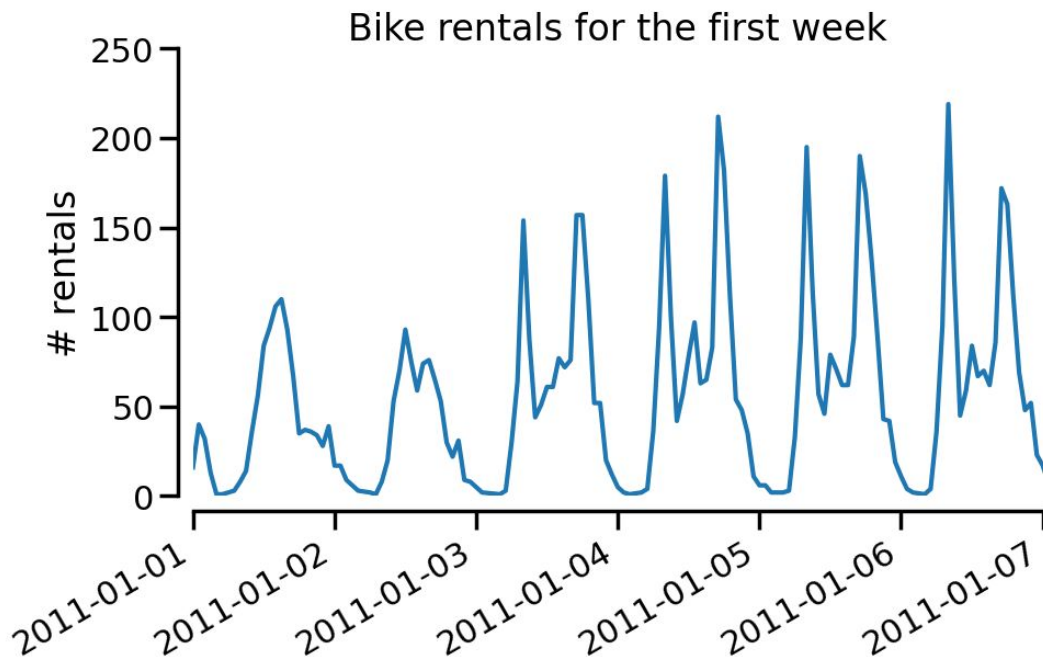
```
from sklearn.datasets import fetch_openml

bike_sharing = fetch_openml(
    "Bike_Sharing_Demand", version=2, as_frame=True, parser="pandas",
)
bike_sharing.frame
```

	season	year	month	hour	holiday	weekday	workingday	weather	temp	feel_temp	humidity	windspeed	count
dteday													
2011-01-01 00:00:00	spring	0	1	0	False	6	False	clear	9.84	14.395	0.81	0.0000	16
2011-01-01 01:00:00	spring	0	1	1	False	6	False	clear	9.02	13.635	0.80	0.0000	40
2011-01-01 02:00:00	spring	0	1	2	False	6	False	clear	9.02	13.635	0.80	0.0000	32
2011-01-01 03:00:00	spring	0	1	3	False	6	False	clear	9.84	14.395	0.75	0.0000	13
2011-01-01 04:00:00	spring	0	1	4	False	6	False	clear	9.84	14.395	0.75	0.0000	1

Tackling time-series as a regression problem

The bike sharing dataset

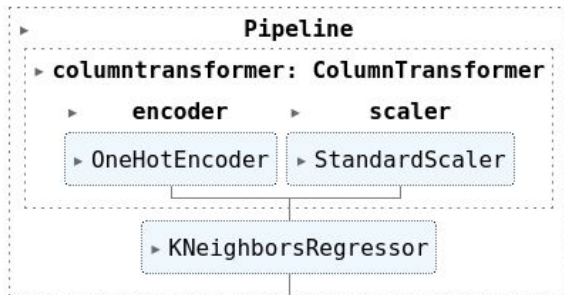


Tackling time-series as a regression problem

Our first (overfitting) baseline model

```
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.pipeline import make_pipeline
from sklearn.neighbors import KNeighborsRegressor

categorical_columns = ["season", "holiday", "workingday", "weather",]
numerical_columns = [
    "year", "month", "hour", "weekday", "temp", "feel_temp", "humidity", "windspeed"
]
preprocessing = ColumnTransformer(transformers=[
    ("encoder", OneHotEncoder(), categorical_columns),
    ("scaler", StandardScaler(), numerical_columns),
])
model = make_pipeline(preprocessing, KNeighborsRegressor(n_neighbors=1))
model
```



Tackling time-series as a regression problem

Evaluation with non-i.i.d data

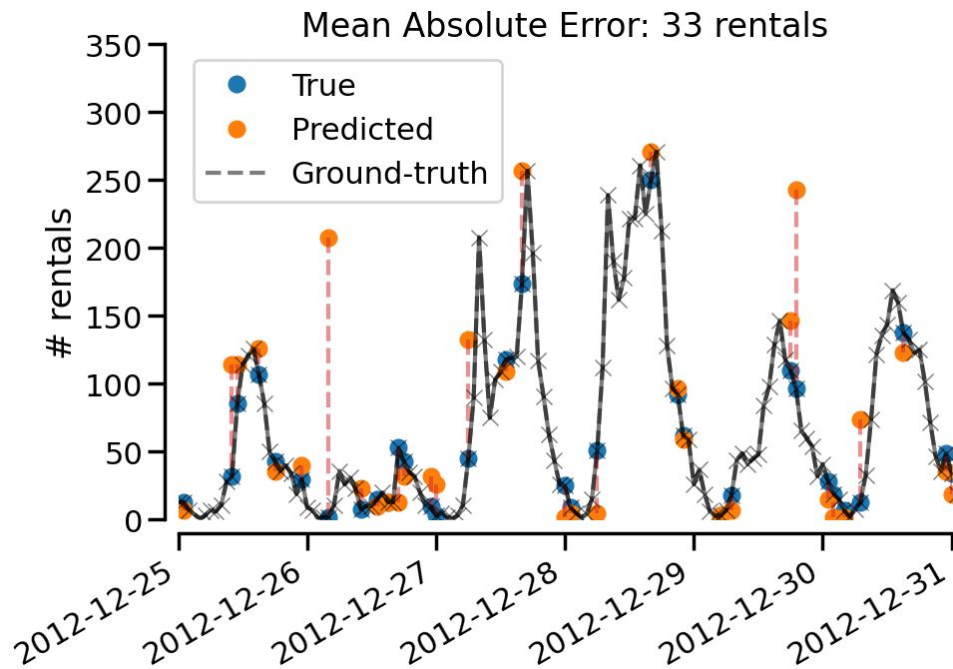
```
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(
    X, y, shuffle=True, random_state=42, test_size=0.2
)
```

```
from sklearn.metrics import mean_absolute_error

model.fit(X_train, y_train)
y_pred = model.predict(X_test)
print(f"MAE: {mean_absolute_error(y_test, y_pred):.0f} rentals")
```

MAE: 76 rentals



Tackling time-series as a regression problem

Evaluation with non-i.i.d data

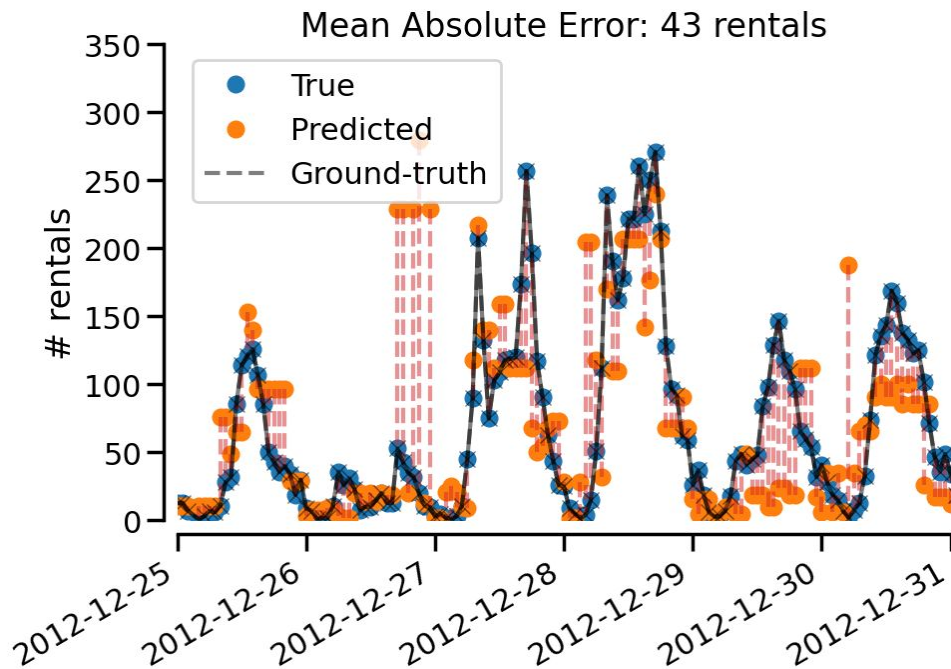
```
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(  
    X, y, shuffle=False, random_state=42, test_size=0.2  
)
```

```
from sklearn.metrics import mean_absolute_error
```

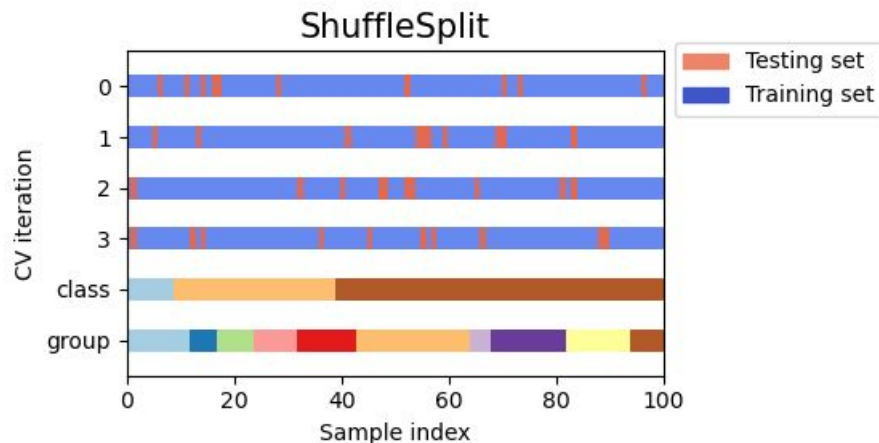
```
model.fit(X_train, y_train)  
y_pred = model.predict(X_test)  
print(f"MAE: {mean_absolute_error(y_test, y_pred):.0f} rentals")
```

MAE: 119 rentals



Tackling time-series as a regression problem

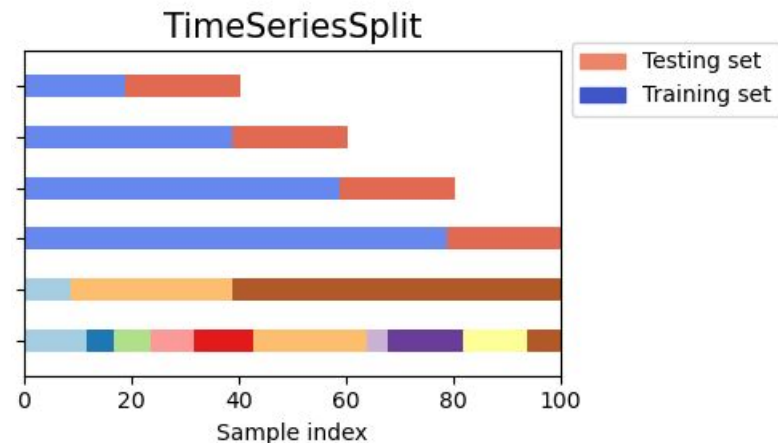
Evaluation with non-i.i.d data



```
from sklearn.model_selection import cross_validate, ShuffleSplit

cv = ShuffleSplit()
results = cross_validate(model, X, y, cv=cv, scoring="neg_mean_absolute_error")
print(f"MAE: {-results['test_score'].mean():.0f} rentals")
```

MAE: 75 rentals



```
from sklearn.model_selection import cross_validate, TimeSeriesSplit

cv = TimeSeriesSplit()
results = cross_validate(model, X, y, cv=cv, scoring="neg_mean_absolute_error")
print(f"MAE: {-results['test_score'].mean():.0f} rentals")
```

MAE: 116 rentals

Tackling time-series as a regression problem

Evaluation with non-i.i.d data

```
from sklearn.model_selection import cross_validate, ShuffleSplit  
  
cv = ShuffleSplit()  
results = cross_validate(model, X, y, cv=cv)  
results["test_score"]
```

```
array([0.51609683, 0.5774131 , 0.53495167, 0.55402703, 0.55465608,  
       0.55297954, 0.591685  , 0.52655122, 0.53399944, 0.5687863 ])
```

```
from sklearn.model_selection import cross_validate, TimeSeriesSplit  
  
cv = TimeSeriesSplit()  
results = cross_validate(model, X, y, cv=cv)  
results["test_score"]
```

```
array([-0.06406429,  0.01166642,  0.03483939,  0.12561052,  0.3277383 ])
```

$$R^2(y, \hat{y}) = 1 - \frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{\sum_{i=1}^n (y_i - \bar{y})^2}$$

Tackling time-series as a regression problem

Evaluation with non-i.i.d data

```
import numpy as np
from sklearn.model_selection import cross_validate
from sklearn.metrics import mean_pinball_loss
from sklearn.metrics import mean_absolute_error
from sklearn.metrics import mean_absolute_error

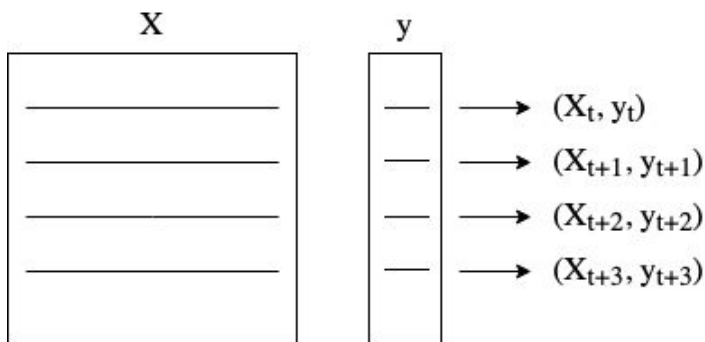
def evaluate(model, X, y, cv):
    def score_func(estimator, X, y):
        y_pred = estimator.predict(X)
        return {
            "mean_absolute_error": mean_absolute_error(y, y_pred),
            "mean_pinball_05_loss": mean_pinball_loss(y, y_pred, alpha=0.05),
            "mean_pinball_50_loss": mean_pinball_loss(y, y_pred, alpha=0.50),
            "mean_pinball_95_loss": mean_pinball_loss(y, y_pred, alpha=0.95),
        }
    cv_results = cross_validate(model, X, y, cv=cv, scoring=score_func)
    for key, value in cv_results.items():
        if key.startswith("test_"):
            print(f"{key[5:]}: {value.mean():.3f} ± {value.std():.3f}")
```

Forecasting vs. regression

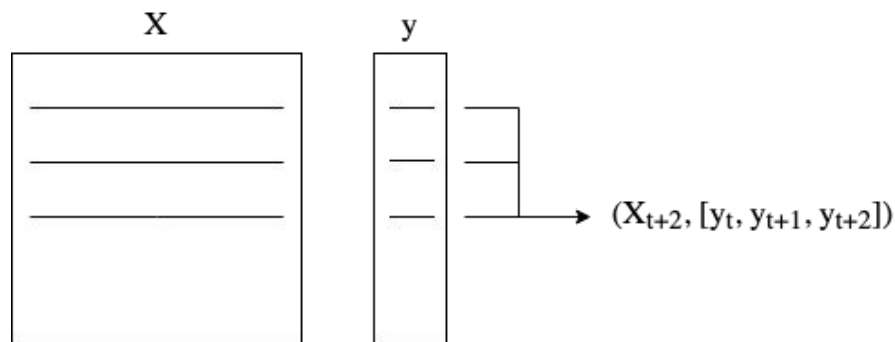
Modeling time dependence

Forecasting is the process of making predictions based on **past** and **present** data [1].

Regression

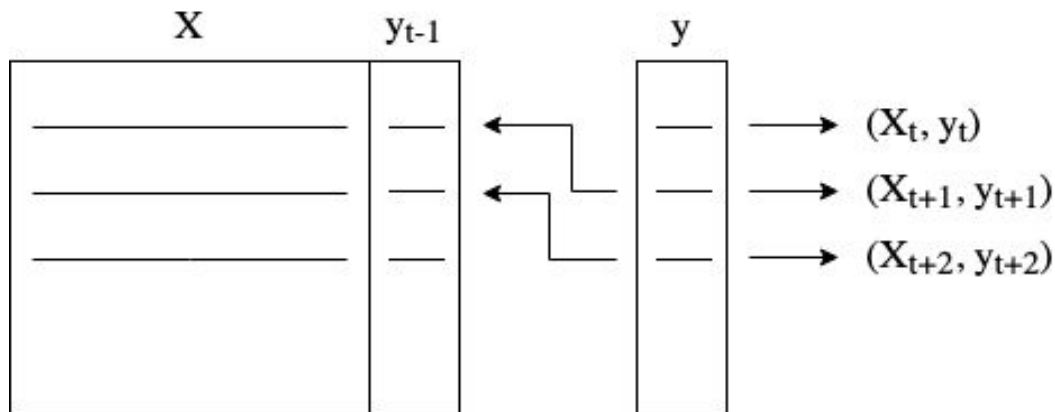


Forecasting



Time-series forecasting via tabularization

Injecting delayed target information in X



Time-series forecasting via tabularization

Delaying target with Pandas

```
lagged_target = pd.concat(  
    [  
        y,  
        y.shift(1).rename("lagged_count_1h"),  
        y.shift(2).rename("lagged_count_2h"),  
        y.shift(3).rename("lagged_count_3h"),  
        y.shift(24).rename("lagged_count_1d"),  
        y.shift(24 + 1).rename("lagged_count_1d_1h"),  
        y.shift(7 * 24).rename("lagged_count_7d"),  
        y.shift(7 * 24 + 1).rename("lagged_count_7d_1h"),  
        y.shift(1).rolling(24).mean().rename("lagged_mean_24h"),  
        y.shift(1).rolling(24).max().rename("lagged_max_24h"),  
        y.shift(1).rolling(24).min().rename("lagged_min_24h"),  
        y.shift(1).rolling(7 * 24).mean().rename("lagged_mean_7d"),  
        y.shift(1).rolling(7 * 24).max().rename("lagged_max_7d"),  
        y.shift(1).rolling(7 * 24).min().rename("lagged_min_7d"),  
    ],  
    axis="columns",  
)
```

lagged_target.head()

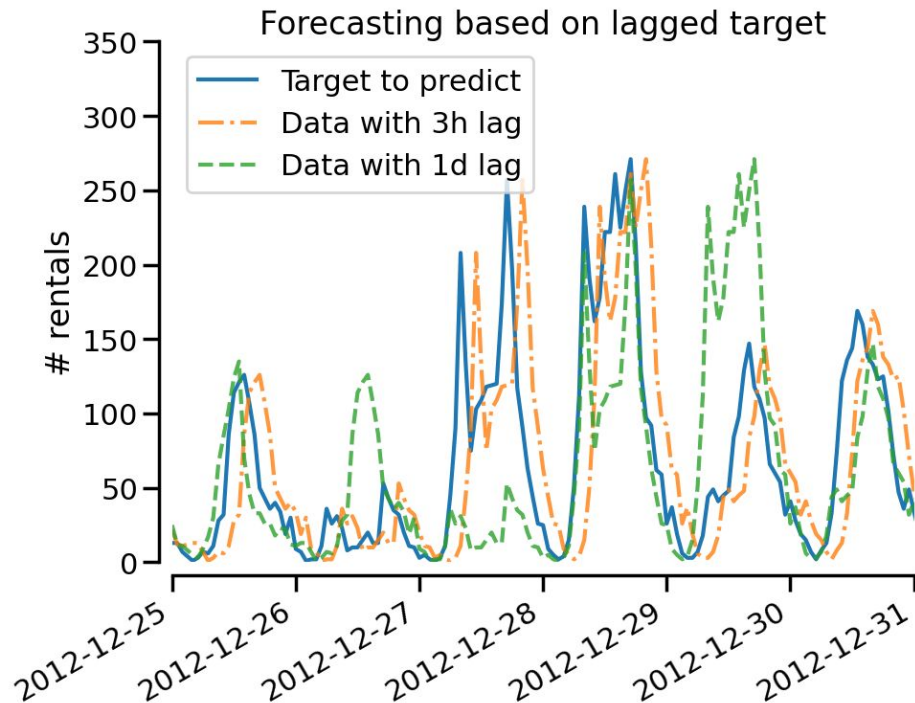
	count	lagged_count_1h	lagged_count_2h	lagged_count_3h	lagged_count_1d	lagged_count_1d_1h	lagged_count_7d
dteday							
2011-01-01 00:00:00	16	NaN	NaN	NaN	NaN	NaN	NaN
2011-01-01 01:00:00	40	16.0	NaN	NaN	NaN	NaN	NaN
2011-01-01 02:00:00	32	40.0	16.0	NaN	NaN	NaN	NaN
2011-01-01 03:00:00	13	32.0	40.0	16.0	NaN	NaN	NaN
2011-01-01 04:00:00	1	13.0	32.0	40.0	NaN	NaN	NaN

lagged_target.tail()

	count	lagged_count_1h	lagged_count_2h	lagged_count_3h	lagged_count_1d	lagged_count_1d_1h	lagged_count_7d
dteday							
2012-12-31 19:00:00	119	122.0	164.0	214.0	102.0	125.0	26.0
2012-12-31 20:00:00	89	119.0	122.0	164.0	72.0	102.0	18.0
2012-12-31 21:00:00	90	89.0	119.0	122.0	47.0	72.0	23.0
2012-12-31 22:00:00	61	90.0	89.0	119.0	36.0	47.0	22.0
2012-12-31 23:00:00	49	61.0	90.0	89.0	49.0	36.0	12.0

Time-series forecasting via tabularization

Delaying target with Pandas



Time-series forecasting via tabularization

Linear models

```
from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import Ridge

model = make_pipeline(StandardScaler(), Ridge(alpha=1e-4))
evaluate(model, X, y)
```

```
mean_absolute_error: 43.204 ± 3.057
mean_pinball_05_loss: 21.638 ± 2.077
mean_pinball_50_loss: 21.602 ± 1.529
mean_pinball_95_loss: 21.566 ± 1.393
```

```
from sklearn.linear_model import PoissonRegressor

model = make_pipeline(StandardScaler(), PoissonRegressor(alpha=1e-4))
evaluate(model, X, y)
```

```
mean_absolute_error: 101.267 ± 12.199
mean_pinball_05_loss: 65.596 ± 11.262
mean_pinball_50_loss: 50.633 ± 6.099
mean_pinball_95_loss: 35.670 ± 5.864
```

Time-series forecasting via tabularization

Gradient boosting regression trees

```
from sklearn.ensemble import HistGradientBoostingRegressor
```

```
gbrt_mse = HistGradientBoostingRegressor(loss="squared_error")  
evaluate(gbrt_mse, X, y)
```

```
mean_absolute_error: 39.088 ± 2.268  
mean_pinball_05_loss: 17.700 ± 1.275  
mean_pinball_50_loss: 19.544 ± 1.134  
mean_pinball_95_loss: 21.388 ± 2.363
```

```
gbrt_poisson = HistGradientBoostingRegressor(loss="poisson")  
evaluate(gbrt_poisson, X, y)
```

```
mean_absolute_error: 39.307 ± 2.808  
mean_pinball_05_loss: 16.669 ± 1.541  
mean_pinball_50_loss: 19.653 ± 1.404  
mean_pinball_95_loss: 22.638 ± 2.983
```

Time-series forecasting via tabularization

Linear models with feature engineering

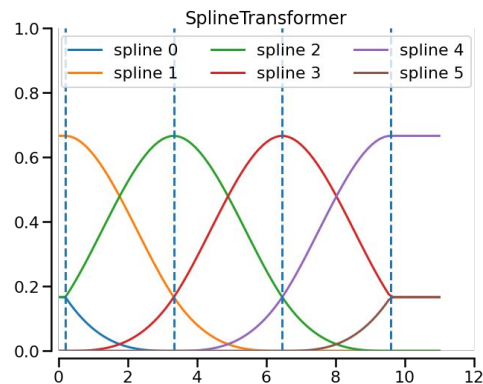
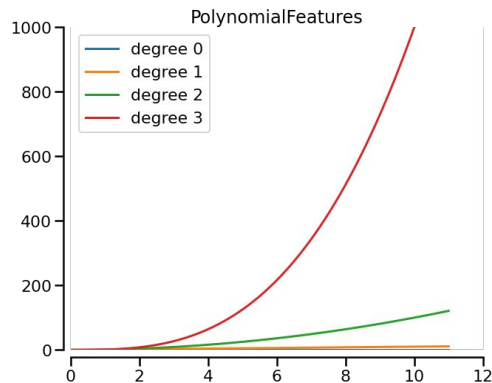
```
X.shape
```

```
(17210, 13)
```

```
from sklearn.preprocessing import SplineTransformer
```

```
SplineTransformer(n_knots=10, degree=3).fit_transform(X).shape
```

```
(17210, 156)
```



Time-series forecasting via tabularization

Linear models with feature engineering

```
import warnings
from sklearn.exceptions import ConvergenceWarning
from sklearn.pipeline import make_pipeline

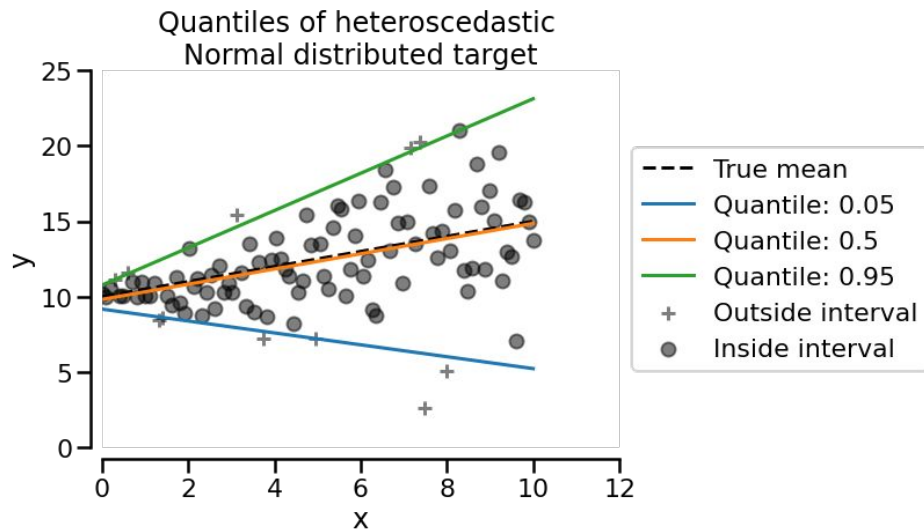
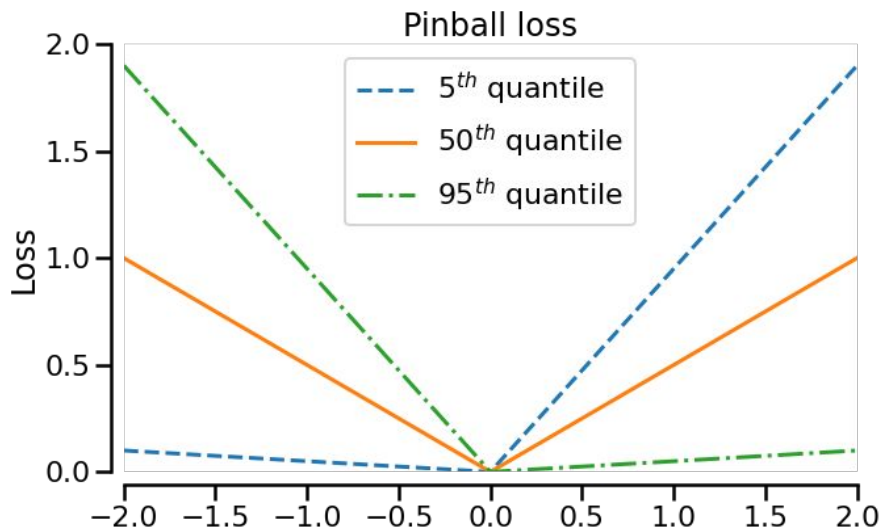
pipeline = make_pipeline(
    SplineTransformer(n_knots=12, degree=3),
    PoissonRegressor(alpha=1e-6, max_iter=300),
)

with warnings.catch_warnings():
    warnings.simplefilter("ignore", ConvergenceWarning)
    evaluate(pipeline, X, y)
```

```
mean_absolute_error: 54.486 ± 17.952
mean_pinball_05_loss: 16.356 ± 2.416
mean_pinball_50_loss: 27.243 ± 8.976
mean_pinball_95_loss: 38.130 ± 19.391
```

Modeling predictive uncertainty

Estimating quantiles



Modeling predictive uncertainty

Gradient boosting regression trees with quantile loss

```
gbrt_percentile_05 = HistGradientBoostingRegressor(loss="quantile", quantile=0.05)  
evaluate(gbrt_percentile_05, X, y)
```

```
mean_absolute_error: 92.476 ± 16.235  
mean_pinball_05_loss: 5.874 ± 0.925  
mean_pinball_50_loss: 46.238 ± 8.117  
mean_pinball_95_loss: 86.602 ± 15.310
```

```
gbrt_median = HistGradientBoostingRegressor(loss="quantile", quantile=0.5)  
evaluate(gbrt_median, X, y)
```

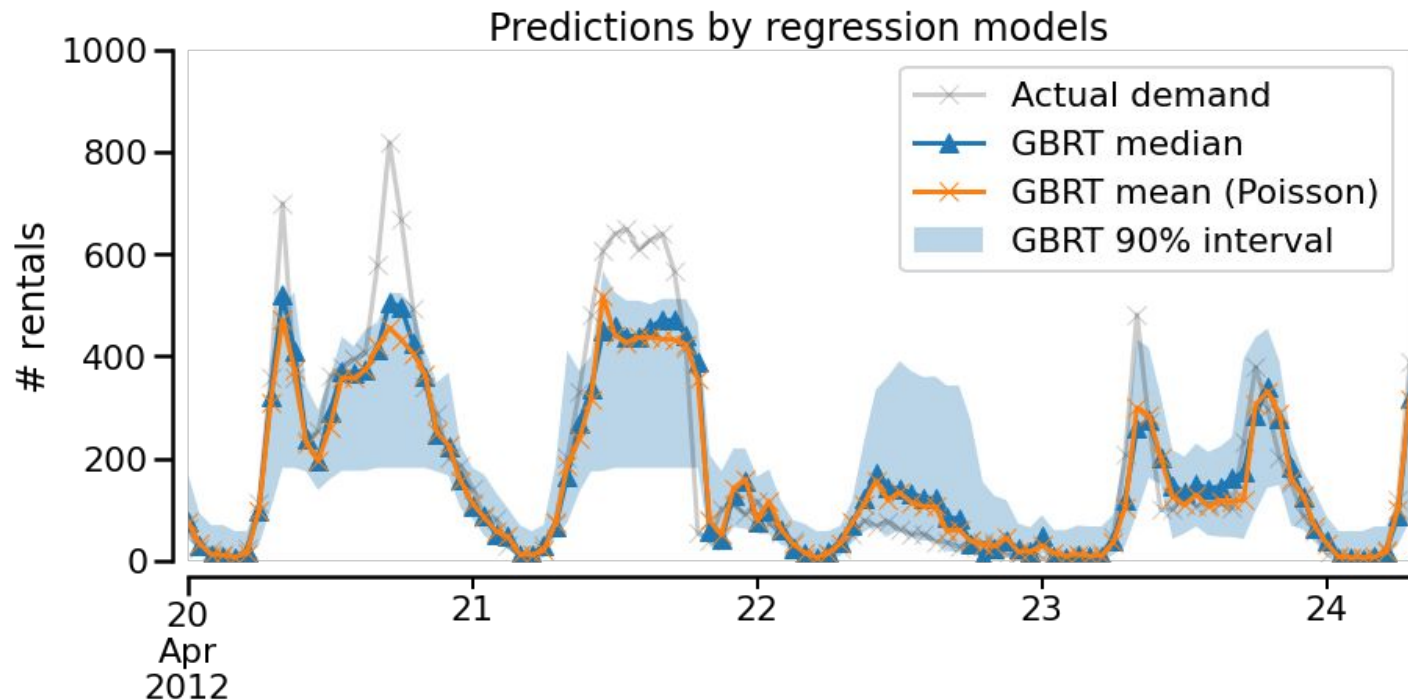
```
mean_absolute_error: 39.854 ± 3.167  
mean_pinball_05_loss: 17.147 ± 1.067  
mean_pinball_50_loss: 19.927 ± 1.584  
mean_pinball_95_loss: 22.706 ± 3.131
```

```
gbrt_percentile_95 = HistGradientBoostingRegressor(loss="quantile", quantile=0.95)  
evaluate(gbrt_percentile_95, X, y)
```

```
mean_absolute_error: 72.009 ± 6.143  
mean_pinball_05_loss: 62.901 ± 7.443  
mean_pinball_50_loss: 36.005 ± 3.071  
mean_pinball_95_loss: 9.109 ± 1.305
```

Modeling predictive uncertainty

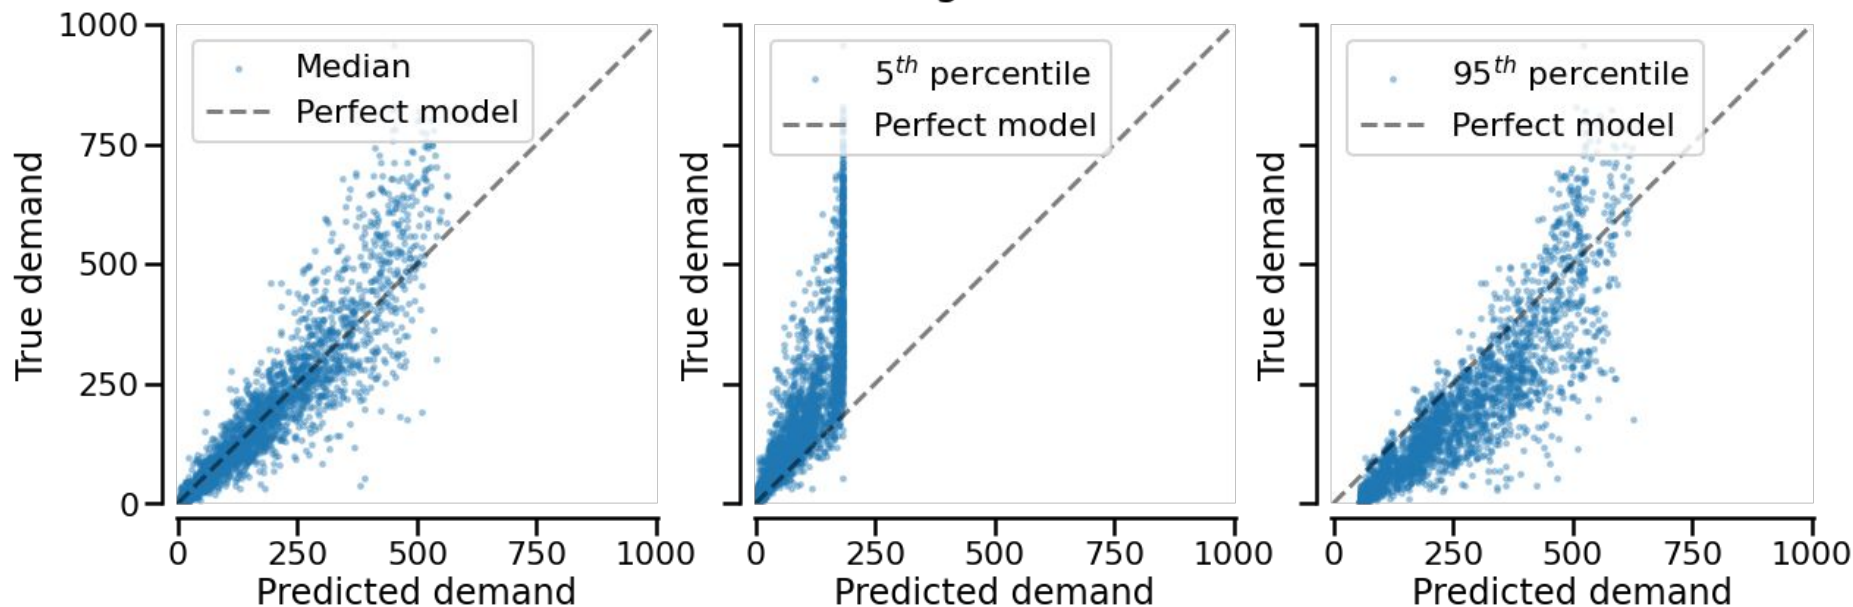
Qualitative evaluation of quantile calibration



Modeling predictive uncertainty

Qualitative evaluation of quantile calibration

Non-linear regression models



Modeling predictive uncertainty

Quantile calibration through effective coverage

```
(median_predictions > y_test).mean()
```

```
0.5246666666666666
```

```
(percentile_5_predictions > y_test).mean()
```

```
0.081
```

```
(percentile_95_predictions > y_test).mean()
```

```
0.8823333333333333
```

```
np.logical_and(  
    percentile_5_predictions < y_test,  
    percentile_95_predictions > y_test,  
) .mean()
```

```
0.8013333333333333
```

Modeling predictive uncertainty

Causes of miscalibrated models

In theory

The pinball loss is guaranteed to be minimized by models that estimate the quantiles perfectly.

In practice

This is an asymptotic property with access to an infinite number of data points.

With a finite set of samples, models with the same pinball loss can **trade calibration for ranking power**

Modeling predictive uncertainty

Can we do better?

```
from numpy import quantile
```

```
fixed_05_quantile, fixed_95_quantile = np.quantile(y_train, [0.05, 0.95])  
fixed_05_quantile, fixed_95_quantile
```

```
(4.0, 423.0)
```

```
np.logical_and(  
    fixed_05_quantile < y_test, fixed_95_quantile > y_test,  
) .mean()
```

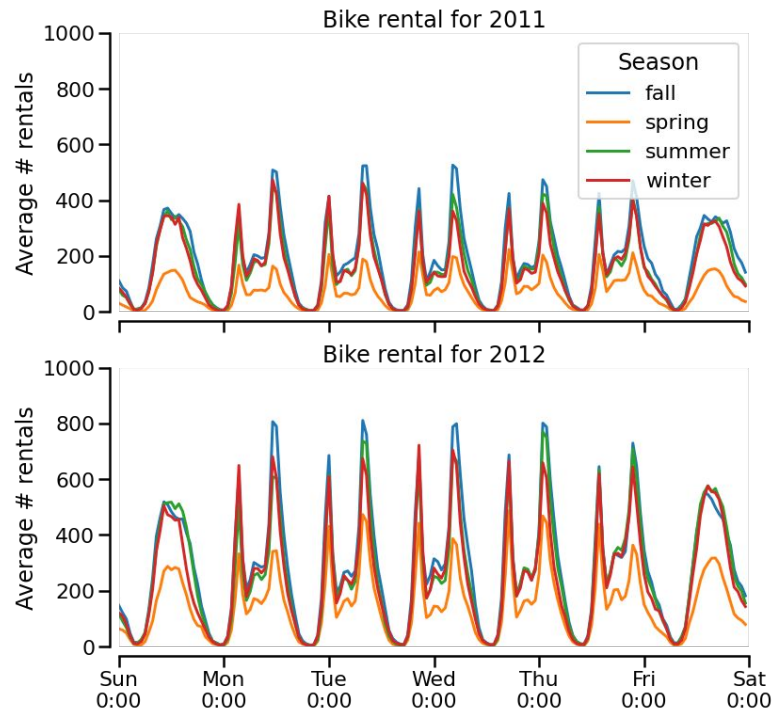
```
0.831
```

```
mean_pinball_loss(  
    y_test, np.full_like(y_test, fill_value=fixed_05_quantile), alpha=0.05  
)
```

```
8.624333333333333
```

```
mean_pinball_loss(  
    y_test, np.full_like(y_test, fill_value=fixed_95_quantile), alpha=0.95  
)
```

```
26.398333333333344
```



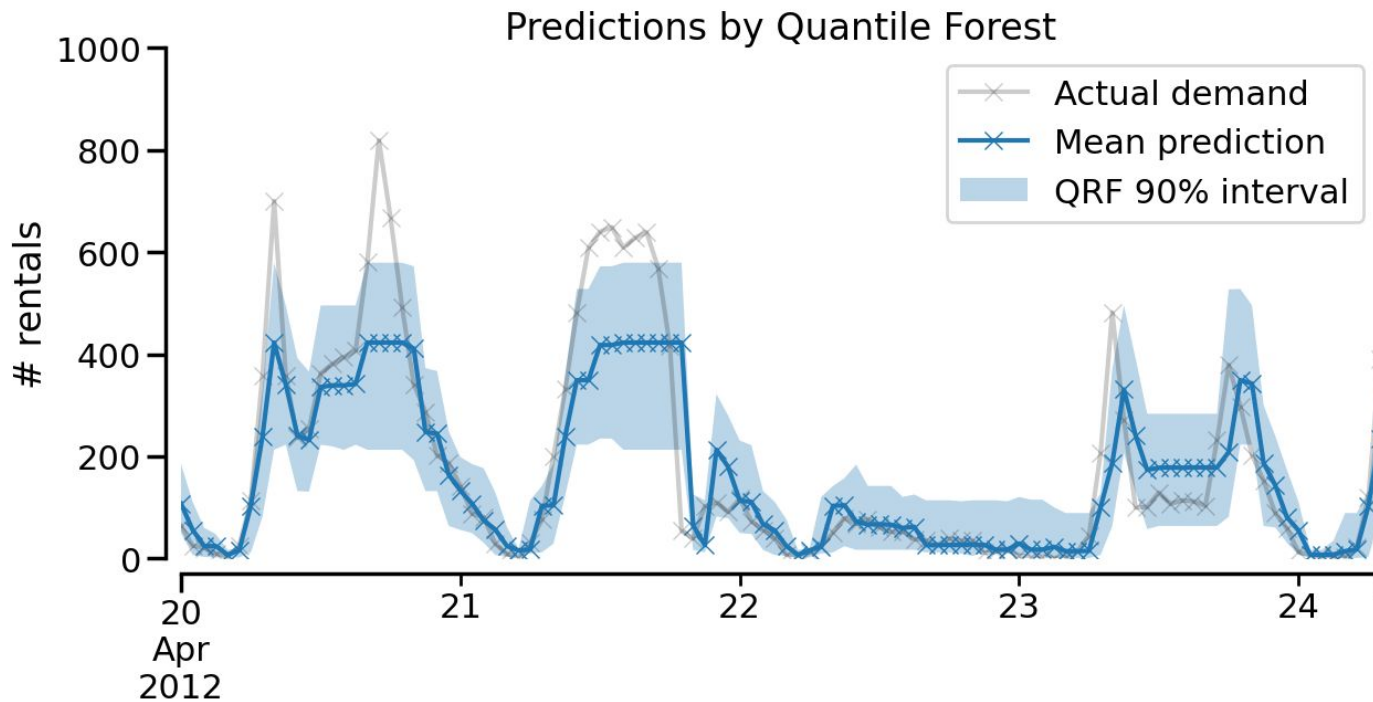
Modeling predictive uncertainty

Can we do better?

- Collecting more data-points (in case the model is overfitting)
- Better tuning of the model hyper-parameters
- Engineering more predictive features from the same data
- Integrate features that will reduce the uncertainty
- Try other kinds of quantile regression models

Modeling predictive uncertainty

Uncertainty prediction with quantile forest



Modeling predictive uncertainty

Uncertainty prediction with quantile forest

```
(percentile_5_predictions > y_test).mean()
```

```
0.052
```

```
(percentile_95_predictions > y_test).mean()
```

```
0.9063333333333333
```

```
np.logical_and(  
    percentile_5_predictions < y_test,  
    percentile_95_predictions > y_test,  
) .mean()
```

```
0.8393333333333334
```

Modeling predictive uncertainty

Uncertainty prediction MAPIE

Uncertainty prediction based on conformal prediction methods and estimates both **aleatoric** and **epistemic** uncertainty at the same time:

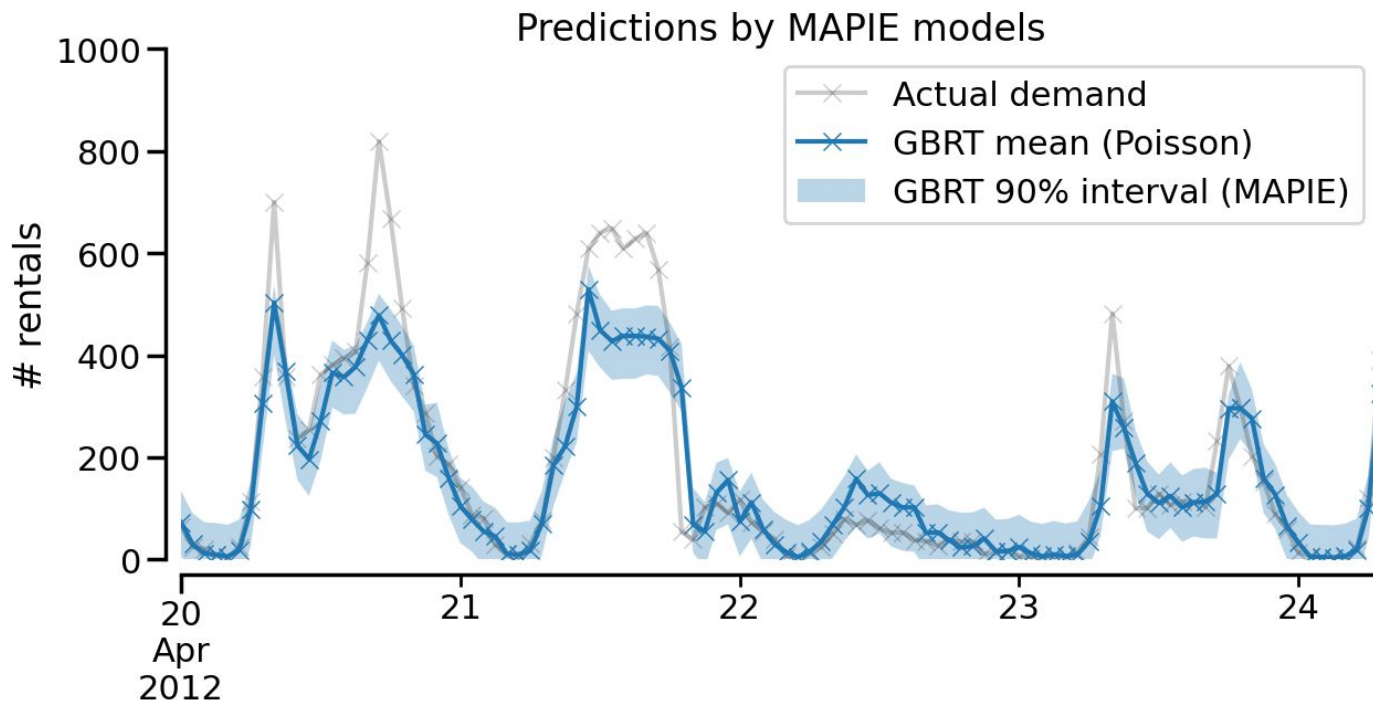
- aleatoric uncertainty due to limited information in the input data;
- epistemic uncertainty due to a limited number of samples.

```
from mapie.regression import MapieRegressor

gbrt_mean_poisson_mapie = MapieRegressor(
    HistGradientBoostingRegressor(loss="poisson"), cv=5
)
gbrt_mean_poisson_mapie.fit(X_train, y_train)
mean_predictions, predictions_90_pi = gbrt_mean_poisson_mapie.predict(X_test, alpha=0.1)
```

Modeling predictive uncertainty

Uncertainty prediction MAPIE



Modeling predictive uncertainty

Uncertainty prediction MAPIE

```
(predictions_90_pi_low > y_test).mean()
```

```
0.052
```

```
(predictions_90_pi_high > y_test).mean()
```

```
0.9063333333333333
```

```
np.logical_and(  
    predictions_90_pi_low < y_test,  
    predictions_90_pi_high > y_test,  
) .mean()
```

```
0.842
```

Beyond scikit-learn

Limitations of the current approach

- Feature engineering not integrated within the cross-validation
- Predicting a single step or in other words the next hour



Beyond scikit-learn

Example of sktime

```
from sktime.forecasting.model_selection import temporal_train_test_split

y_train, y_test = temporal_train_test_split(y, test_size=7 * 24)
```

```
from sktime.forecasting.base import ForecastingHorizon

fh = ForecastingHorizon(y_test.index, is_relative=False)
fh
```

```
ForecastingHorizon(['2012-12-25 00:00:00', '2012-12-25 01:00:00',
                    '2012-12-25 02:00:00', '2012-12-25 03:00:00',
                    '2012-12-25 04:00:00', '2012-12-25 05:00:00',
                    '2012-12-25 06:00:00', '2012-12-25 07:00:00',
                    '2012-12-25 08:00:00', '2012-12-25 09:00:00',
                    ...,
                    '2012-12-31 14:00:00', '2012-12-31 15:00:00',
                    '2012-12-31 16:00:00', '2012-12-31 17:00:00',
                    '2012-12-31 18:00:00', '2012-12-31 19:00:00',
                    '2012-12-31 20:00:00', '2012-12-31 21:00:00',
                    '2012-12-31 22:00:00', '2012-12-31 23:00:00'],
                    dtype='datetime64[ns]', name='dteday', length=168, freq='H', is_relative=
```

Beyond scikit-learn

Example of sktime

```
from sktime.forecasting.compose import make_reduction

regressor = HistGradientBoostingRegressor(
    loss="poisson", max_leaf_nodes=64, max_iter=300
)
forecaster = make_reduction(regressor, window_length=7 * 24, strategy="recursive")
forecaster.fit(y_train)
```

► **RecursiveTabularRegressionForecaster**

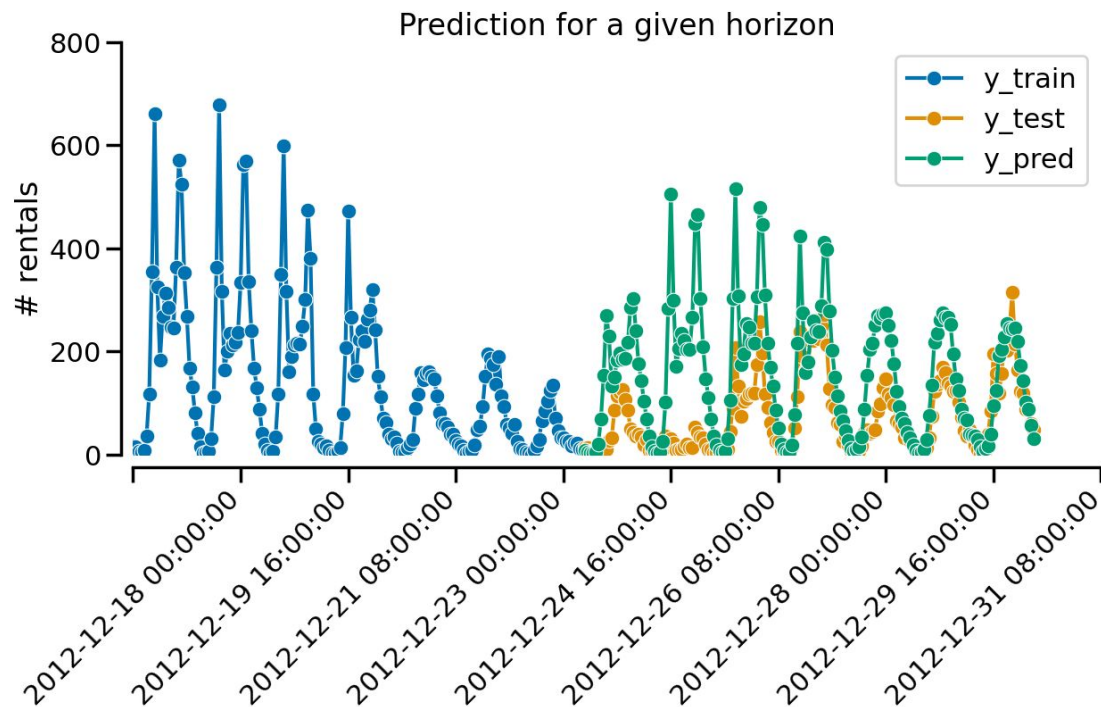
► **estimator: HistGradientBoostingRegressor**

► HistGradientBoostingRegressor

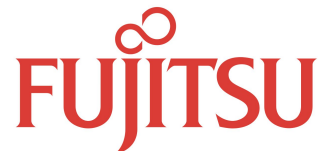
```
y_pred = forecaster.predict(fh=fh)
```

Beyond scikit-learn

Example of sktime



scikit-learn @ Inria foundation partners:



Hugging Face



data
iku



Tutorial material

GitHub repository

ogrisel / euroscipy-2022-time-series Public

Watch 1

Fork 6

Star 29

<> Code Issues Pull requests Actions Projects Security Insights

main 1 branch 0 tags

Go to file

Add file

Code

About





Tutorial on time-series forecasting with scikit-learn

MIT license

29 stars

1 watching

6 forks

 ogrisel	Include QRF	14c9ae5 on Sep 7	🕒 12 commits
	.gitignore	Initial commit	3 months ago
	LICENSE	Initial commit	3 months ago
	plot_time_series_feature_engineer...	Include QRF	2 months ago

<https://github.com/ogrisel/euroscipy-2022-time-series>